

# Linear Regression with `scikit-learn`

# Linear Regression



## scikit-learn: Using Linear Algebra

```
from sklearn.linear_model import LinearRegression  
  
lin_reg = LinearRegression()  
lin_reg.fit(X, y)  
lin_reg.coef_, lin_reg.intercept_
```

Note: this is not a learning algorithm.

# Optimizers

1. (Batch) Gradient Descent
2. Stochastic Gradient Descent
3. Mini-Batch Gradient Descent

# scikit-learn: Stochastic Gradient Descent

```
from sklearn.linear_model import SGDRegressor

sgd_reg = SGDRegressor()
sgd_reg.fit(X, y)
sgd_reg.coef_, sgd_reg.intercept_
```

## scikit-learn: SGD Hyperparameters

```
from sklearn.linear_model import SGDRegressor

sgd_reg = SGDRegressor(
    max_iter=100000,
    n_iter_no_change=10,
    tol=1e-4,
    learning_rate='adaptive',
)
sgd_reg.fit(X, y)
sgd_reg.coef_, sgd_reg.intercept_
```

## scikit-learn: SGD `partial_fit`

```
from sklearn.linear_model import SGDRegressor

sgd_reg = SGDRegressor()
sgd_reg.partial_fit(X_1, y_1)
sgd_reg.partial_fit(X_2, y_2)
...
sgd_reg.partial_fit(X_n, y_n)
sgd_reg.coef_, sgd_reg.intercept_
```

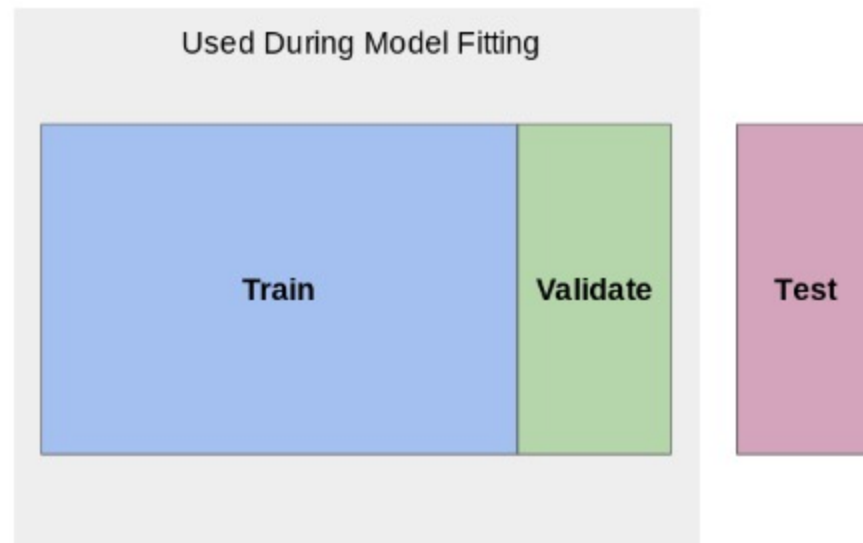
# Loss

## Mean Squared Error

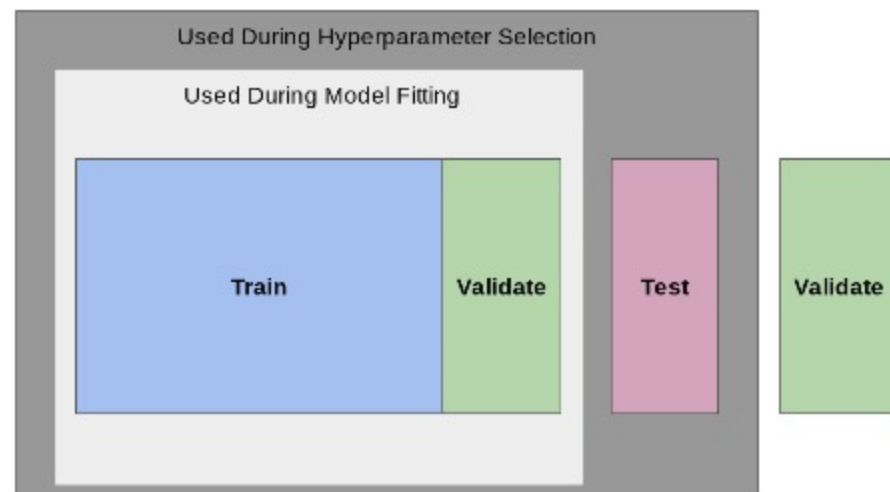
$$MSE = \frac{1}{n} \sum_{n=1}^n (y_i - \hat{y}_i)^2$$



# Train/Validate, Test



# Train/Validate, Test, Validate



**Your Turn**